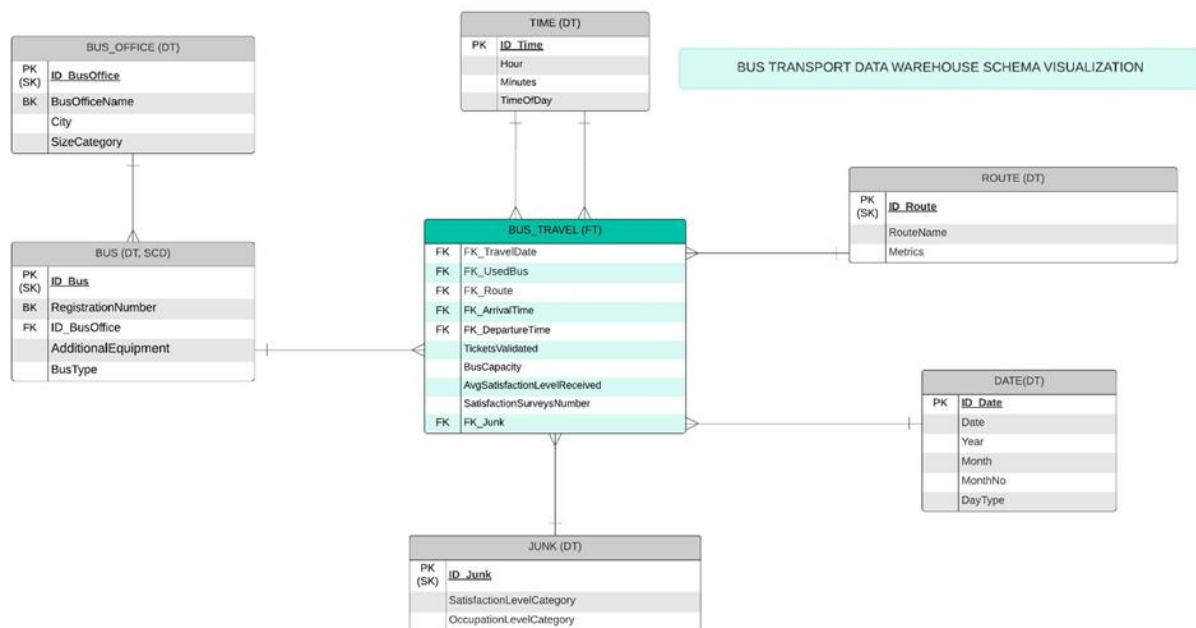


# Bus transport Warehouse Design in Hive

## Scenarios and explanation of the decisions made



## Competency questions and scenarios

### 1. What are the weekday routes with the highest number of tickets validated?(comparing two months to each other)

**Scenario:** A transport manager is analyzing weekday bus operations to identify the most profitable or popular routes, in order to allocate more buses for these routes.

**Partitioning** by **day\_type** (weekday) and **month\_no** allows the query to efficiently filter for weekdays in given months.

**Bucketing** on **route\_id** helps speed up the join between travel and route tables.

```

SELECT r.route_id, SUM(CASE WHEN d.month_no = 4 THEN
t.tickets_validated ELSE 0 END) AS total_tickets_april,
SUM(CASE WHEN d.month_no = 5 THEN t.tickets_validated ELSE 0
END) AS total_tickets_may, (SUM(CASE WHEN d.month_no = 5 THEN
t.tickets_validated ELSE 0 END)-SUM(CASE WHEN d.month_no = 4
THEN t.tickets_validated ELSE 0 END)) AS
tickets_difference_to_prev_month
FROM travel t
JOIN route r ON t.route_id = r.route_id
JOIN date_dim d ON t.travel_date = d.date_id
WHERE d.month_no IN (4, 5) AND d.day_type= 'weekday'
GROUP BY r.route_id
ORDER BY total_tickets_april DESC
LIMIT 10;
  
```

r.route_id	total_tickets_april	total_tickets_may	tickets_difference_to_prev_month
86	1243	1011	-232
80	1180	1001	-179
73	1153	1141	-12
32	1147	1060	-87
40	1145	1096	-49
35	1141	1049	-92
84	1139	1045	-94
30	1129	1051	-78
66	1125	1067	-58
4	1095	1175	80

(case clause used in query: <https://stackoverflow.com/questions/31489073/hive-else-in-sum> )

## 2. Which bus type has the highest average satisfaction level?

**Scenario:** The management team wants to understand which type of buses provide the best customer experience to decide on future bus purchases.

**Partitioning** by `bus_type` in the bus table allows efficient retrieval of bus types (on `bus_type` we use static partitioning which ensures that each bus type has its own dedicated partition)

**Bucketing** by `bus_id` helps speed up joins between the travel and bus tables.

```
SELECT b.bus_type, AVG(t.avg_satisfaction_level_received) AS
avg_satisfaction
FROM travel t
JOIN bus b ON t.bus_id = b.bus_id
GROUP BY b.bus_type
ORDER BY avg_satisfaction DESC
LIMIT 1;
```

b.bus_type	avg_satisfaction
low floor	3.7620130793289737

## 3. What is the number of validated tickets for each bus comparing across weekdays and weekends in May?

**Scenario:** A data analyst is looking at bus usage in May to see which buses are used more on weekdays than at weekends to spot which buses have higher demand on weekdays versus weekends.

**Partitioning** on by `month_no` and `day_type` for faster access only the relevant data for May and weekdays/weekends.

**Bucketing** on `bus_id` in the travel table helps improve the efficiency of operations involving `bus_id`.

```

SELECT t.bus_id, d.day_type, SUM(t.tickets_validated) AS
total_validated_tickets
FROM travel t
JOIN date_dim d ON t.travel_date = d.date_id
WHERE d.month_no = 5 AND d.day_type IN ('weekday', 'weekend')
GROUP BY t.bus_id, d.day_type
ORDER BY t.bus_id, d.day_type;

```

t.bus_id	d.day_type	total_validated_tickets
1	weekday	307
1	weekend	84
2	weekday	568
2	weekend	203
3	weekday	693
3	weekend	331
4	weekday	439
4	weekend	150
5	weekday	494
5	weekend	198
6	weekday	703
6	weekend	223
7	weekday	457
7	weekend	89
8	weekday	643
8	weekend	181
9	weekday	755
9	weekend	207
10	weekday	462
10	weekend	160
11	weekday	211
11	weekend	101
12	weekday	449
12	weekend	267
13	weekday	550
13	weekend	219
14	weekday	277

#### 4. What was the busiest time of the day for bus routes on weekends in April?

**Scenario:** The operations team is making changes to the weekend bus schedule in April to make better use of resources and reduce crowds.

**Partitioning** by `day_type` (weekend) allows fast filtering for weekend records.

Partitioning by `month_no` helps quickly focus on April data.

**Bucketing** on `route_id` helps optimize joins between travel and route tables.

```

SELECT td.time_of_day, COUNT(*) AS trip_num
FROM travel t
JOIN date_dim d ON t.travel_date = d.date_id
JOIN time_dim td ON t.departure_time = td.time_id
WHERE month_no = 4 AND d.day_type = 'weekend'
GROUP BY td.time_of_day
ORDER BY trip_num DESC;

```

td.time_of_day	trip_num
Morning	856
Afternoon	784
Evening	632
Night	232

##### 5. How many satisfaction surveys were received for a given route comparing two months (April and May)?

**Scenario:** The quality assurance team wants to check customer engagement by comparing the number of satisfaction surveys collected for a routes in April and May.

**Partitioning** by `month_no` allows the query to quickly filter records for April and May.

**Bucketing** on `route_id` helps optimize joins with the route table.

```
SELECT r.route_name, SUM(CASE WHEN d.month_no = 4 THEN
t.satisfaction_surveys_number ELSE 0 END) as surveys_april,
SUM(CASE WHEN d.month_no = 5 THEN
t.satisfaction_surveys_number ELSE 0 END) as surveys_may
FROM travel t
JOIN route r ON t.route_id = r.route_id
JOIN date_dim d ON t.travel_date = d.date_id
GROUP BY r.route_name
ORDER BY surveys_april DESC
LIMIT 5;
```

r.route_name	surveys_april	surveys_may
Warszawa - Radom	175	160
Warszawa - Torun	171	158
Torun - Gdynia	169	175
Koszalin - Gdansk	167	154
Torun - Olsztyn	163	165

##### 6. How satisfaction level of bus users changed over time?

**Scenario:** The quality assurance team wants to see how (for example) recent changes in service affect customer satisfaction.

**Partitioning** by `month_no` helps quickly retrieve data for specific months.

```
SELECT d.month, AVG(t.avg_satisfaction_level_received) AS
avg_satisfaction_level
FROM travel t
JOIN date_dim d ON t.travel_date = d.date_id
GROUP BY d.month
ORDER BY d.month;
```

d.month	avg_satisfaction_level
April	3.2954275029746727
May	4.072341491938852

## 7. How does the satisfaction level vary based on occupation level categories?

**Scenario:** A quality assurance team evaluates how passenger satisfaction levels change based on bus crowd levels to determine whether overcrowding impacts user satisfaction and identify areas for capacity improvement.

The junk table, which contains occupation categories, is small, so partitioning or bucketing isn't critical here.

```
SELECT j.occupation_level_category,
AVG(avg_satisfaction_level_received) AS avg_satisfaction_level
FROM travel t
JOIN junk j ON t.junk_id=j.junk_id
GROUP BY j.occupation_level_category
ORDER BY avg_satisfaction_level DESC;
```

j.occupation_level_category	avg_satisfaction_level
medium	3.7784625364904314
high	3.7640126907387823
low	3.4103308183401047
very low	3.180597014925373

## 8. What is the breakdown of travel counts by bus type that departures in the morning?

(how many buses departed in the morning until now and what bus type was it)

**Scenario:** The transport planner analyses the number of buses of each type that departures in the morning to identify patterns of use and to ensure a balanced use of the different bus types.

**Partitioning** by [bus\\_type](#) allows efficient filtering of travel records by bus type, partitioning by [time\\_of\\_day](#) (morning) filters morning travel data.

**Bucketing** by [bus\\_id](#) can help improve performance when joining the travel table with the bus table.

```
SELECT b.bus_type, COUNT(*) travels_num
FROM travel t
JOIN bus b ON t.bus_id=b.bus_id
JOIN time_dim tdim ON t.departure_time=tdim.time_id
WHERE tdim.time_of_day = 'Morning'
GROUP BY b.bus_type, tdim.time_of_day
```

```
ORDER BY travels_num DESC;
```

b.bus_type	travels_num
standard	5938
minibus	2523
low floor	2325

## 9. Which are the routes with the highest distance value?

**Scenario:** Identifying long-distance routes for subsequent assessment of their viability, fuel consumption and resource requirements.

**Mapping** the `distance_km` field inside `metrics['distance_km']`. Because bucketing by `route_id` ensures that the data related to each route is grouped together, querying the `metrics['distance_km']` map field becomes more efficient.

```
SELECT route_id, route_name, metrics['distance_km'] AS
distance
FROM route
ORDER BY distance DESC
LIMIT 10;
```

route_id	route_name	distance
12	Torun - Olsztyn	497
52	Radom - Bydgoszcz	488
59	Warszawa - Torun	487
14	Torun - Warszawa	486
22	Gdansk - Radom	480
83	Szczecin - Elblag	478
29	Gdynia - Szczecin	478
30	Gdynia - Koszalin	468
75	Koszalin - Gdynia	454
50	Elblag - Bydgoszcz	452

## 10. What is the satisfaction level on weekend routes that have the highest avg duration? (in the selected month - April)

**Scenario:** A quality control team looks at how customers feel about long weekend trips.

**Partitioning** by `day_type` (weekend) and `month_no` (April) allows for efficient filtering of month and day type.

**Bucketing** on `route_id` speeds up the join with the route table.

**Mapping** the `duration_min` in `metrics` for each route, and bucketing by `route_id` helps optimize the retrieval of `metrics['duration_min']`

```
SELECT t.route_id, r.route_name,
AVG(r.metrics['duration_min']) as avg_duration,
AVG(t.avg_satisfaction_level_received) AS avg_satisfaction
```

```

FROM travel t
JOIN route r ON t.route_id = r.route_id
JOIN date_dim d ON t.travel_date = d.date_id
WHERE d.month_no = 4 AND d.day_type = 'weekend'
GROUP BY t.route_id, r.route_name
ORDER BY avg_duration DESC
LIMIT 10;

```

t.route_id	r.route_name	avg_duration	avg_satisfaction
12	Torun - Olsztyn	738.0	3.5925925925925926
83	Szczecin - Elblag	724.0	3.8888888888888889
59	Warszawa - Torun	723.0	4.4305555555555555
52	Radom - Bydgoszcz	721.0	5.0694444444444445
14	Torun - Warszawa	715.0	4.1944444444444445
22	Gdansk - Radom	709.0	3.1527777777777777
30	Gdynia - Koszalin	706.0	4.3333333333333333
29	Gdynia - Szczecin	704.0	3.6944444444444446
75	Koszalin - Gdynia	692.0	3.375
50	Elblag - Bydgoszcz	678.0	3.0694444444444446

## **1. Partitioning:**

Used for tables where queries often filter by specific columns like month\_no, day\_type, time\_of\_day, travel\_date, bus\_type. It limits the number of nodes that needs to be scanned when filtering some columns.

Partitions in tables are comparatively equal size (travel\_date in travel table as the schedule of buses does not vary significantly on different days) or the values possible for each field is limited (like month\_no, day\_type, time\_of\_day and bus\_type). Static partitioning is used for bus\_type where we manually decide on number of partitions and to which partition data will be loaded, we need to be sure that we put data into right partition e.g. loading buses of only one type at once! For dynamic partitioning (month\_no, day\_type, time\_of\_day, travel\_date) - partitions are created automatically depending on data we are inserting into, what for travel\_date field adding records while manually giving the date may be more complex process, probably this table will grow large in time.

## **2. Bucketing:**

Bucket by route\_id, bus\_id in travel table to improve performance of aggregations and joins. Partitioning here might lead to some larger and some smaller partitions as we assume that some routes have more travels – we use buckets to decompose dataset based on value from hash function. As the number of travels is around 20k (data of travels is from April) we can also imply that the number of records in the datawarehouse will arise at most to 240k – when storing data for one year, and we have 90 unique routes we decided on splitting data into 50 buckets.

Bucketing also on route\_id in routes table improves joins and searching on this column, we have 90 routes so splitting into 10 buckets is reasonable (on average 9 routes per bucket). In queries the route\_id field (what is unique attribute) will be used often in queries (e.g. searching for specific routes usually combined with travel table) so map-side joins will work faster on bucketing tables.

Bucketing on bus\_id into 10 buckets and date\_id into 30 buckets – for joins with travel table.

No partitioning/bucketing on junk table as it has only 12 records.

## **3. Data Format:**

ORC – used for routes and buses as it support complex datatypes (in routes table map, in bus array) as well as queries that are using aggregations (travel is join-heavy with other tables).

PARQUET – used for date, time, junk tables, frequently queried for filtering/grouping (day\_type, time\_of\_day). Those tables are relatively small tables. Parquet avoids reading unnecessary columns during query execution (columnar storage).



TEXTFILE – used for bus\_office table; as it may be easier to use for external systems. (reading data in Excel, Python or R) No partitioning here due to relatively small table.